Bruce W. Lee

Google Scholar: scholar.google.com/citations?user=a9HZkjMAAAAJ&hl=en
Github: github.com/brucewlee
Website: brucewlee.com
Email: phys.w.s.lee@gmail.com

Education

Bachelor of Applied Science, Computer Science
University of Pennsylvania, Philadelphia, PA, expected May 2026

Preprints & Reports
†: core contrib.

Distillation Robustifies Unlearning
Lee, B. W.†, Foote, A.†, Infanger, A.†, Shor, L.†, Kamath, H.†, ... & Turner, A. M.
Jun 2025

Utility Engineering: Analyzing and Controlling Emergent Value Systems in AIs
Mazeika, M., Yin, X., Tamirisa, R., Lim, J., **Lee, B. W.**, ... & Hendrycks, D.
Feb 2025

HyperCLOVA X Technical Report
Yoo, K. M., Han, J., In, S., Jeon, H., Jeong, J., ..., **Lee, B. W.**, ... & Jung, J.
Apr 2024

Refereed Publications
∗: equal contrib.

Programming Refusal with Conditional Activation Steering
Lee, B. W., Padhi, I., Ramamurthy, K. N., Miehling, E., ..., & Dhurandhar, A.
ICLR 2025 (Spotlight)

Language Models Don't Learn the Physical Manifestation of Language
Lee, B. W., & Lim, J.
ACL 2024

Instruction Tuning with Human Curriculum
Lee, B. W.∗, Cho, H.∗, & Yoo, K. M.
NAACL 2024

Handcrafted Features in Computational Linguistics
Lee, B. W., & Lee, J. H. J.
BEA @ ACL 2023

Linguistic Properties of Truthful Response
Lee, B. W., Arockiaraj, B. F., & Jin, H.
TrustNLP @ ACL 2023

Prompt-based Learning for Text Readability Assessment
Lee, B. W., & Lee, J.
EACL 2023

Pushing on Text Readability Assessment: A Transformer Meets Handcrafted Linguistic Features
Lee, B. W., Jang, Y. S., & Lee, J. H. J.
EMNLP 2021

Improving Text Readability Assessment Model for L2 English Students in Korea
Lee, B. W. & Lee, J. H. J.
NLP-TEA @ AACL 2020

A Low-cost Cryogenic Temperature Measurement System using Arduino Microcontroller
Lee, W. S.
Physics Education, 55(2)

Simplifying the Vacuum Bazooka
Lee, J., **Lee, W. S**., & Shin, E.
Physics Education, 54(3)

Experience

ML Alignment & Theory Scholars — Berkeley, CA
Research Scholar — Jun 2025 – Present

- **Mentor(s):** Tomek Korbak (UK AI Security Institute)
- Studying strategies to defend against misbehaving agents

ML Alignment & Theory Scholars — Remote
Research Scholar — Jan 2025 – Jun 2025

- **Mentor(s):** Alex Cloud & Alex Turner (Google DeepMind)
- Demonstrated that distillation robustifies machine unlearning: showed existing methods only suppress capabilities behaviorally, while distillation removes them more permanently
- Developed UNDO algorithm that trades compute for unlearning robustness, establishing new Pareto frontier and requiring only 0.01% labeled data compared to full retraining
- Conducted extensive experiments on custom Gemma models (100M-300M params), implementing and benchmarking 6+ unlearning methods across multiple domains

Anthropic — Contract, Remote
Auto Alignment Research Trainer — Nov 2024 – Jun 2025

Center for AI Safety — Remote
Research Collaborator — Sep 2024 – Jan 2025

- **Mentor(s):** Mantas Mazeika
- Developed preference elicitation methods that aim to quantify value representations in LLMs
- Wrote asynchronous Python evaluation scripts to assess value coherence and adversarial risk

IBM Research (Trustworthy AI) — Yorktown Heights, NY
Research Intern — May 2024 – Aug 2024

- **Mentor(s):** Inkit Padhi & Karthikeyan N. Ramamurthy
- Proposed Conditional Activation Steering (CAST), enabling context-dependent LLM control without weight updates; work accepted at ICLR 2025 (Spotlight)
- Built IBM's first activation steering library (github.com/IBM/activation-steering), now adopted across multiple IBM research projects

NAVER Cloud (Hyperclova AI) — South Korea
Research Intern — May 2023 – Aug 2023

- **Mentor(s):** Kang Min Yoo
- Proposed Curriculum Instruction Tuning that structures training data by cognitive complexity
- Helped implement synthetic data generation and instruction tuning pipeline for a proprietary LLM

LXPER South Korea
Research Engineer Apr 2020 – Apr 2023

- Led NLP research at an EdTech startup, architecting production-ready BERT variants for lexical analysis, grammatical error correction, and readability assessment
- Set up AWS-based serverless infrastructures to produce APIs, facilitating the complete lifecycle from research to production rollout

Center for Axion and Precision Physics Research / IBS South Korea
Research Scholar May 2019 – Aug 2019

- **Mentor(s):** Andrei Matlashov
- One of two high school students selected for a prestigious summer physics research program for undergraduate/graduate-level students
- Designed a low-cost Arduino-based cryogenic temperature measurement system, which shows a reasonable accuracy for superconducting quantum interference device (SQUID) experiments

Grants

Career Development and Transition Funding Open Philanthropy, 2025

Long-Term Future Fund EA Funds, 2025

Gutmann-Doyle Research Opportunities Fund UPenn, 2025

Cohere For AI Research Grant Cohere, 2024

Khan Family AI for Business Award UPenn, 2024
Top 13 finalist at Penn Venture Lab Startup Challenge
For an open-source LLM evaluation software, founded a non-profit org

Minister of Science and ICT Award Government of South Korea, 2022
Top 10 submission out of 5420 at a Nationwide Startup Competition
For a transformer-based translator software that allows you to choose writing style

Minister of National Defense Award Government of South Korea, 2022
Top 1 submission out of 953 at a MoND Startup Competition
For a translator software that outperformed Google Translate for narrow use cases

Notable Softwares

IBM/Activation-Steering, **80+★**, 90% Contribution
A popular implementation of activation steering
github.com/IBM/activation-steering

LFTK, **100+★**, 100% Contribution
A multilingual, refactorized version of LingFeat. Cited and used internationally
github.com/brucewlee/lftk

LingFeat, **100+★**, 100% Contribution
A Python library that calculates 255 linguistic features from a text
github.com/brucewlee/lingfeat

Last updated: July 2025.